

# Gesture Recognition

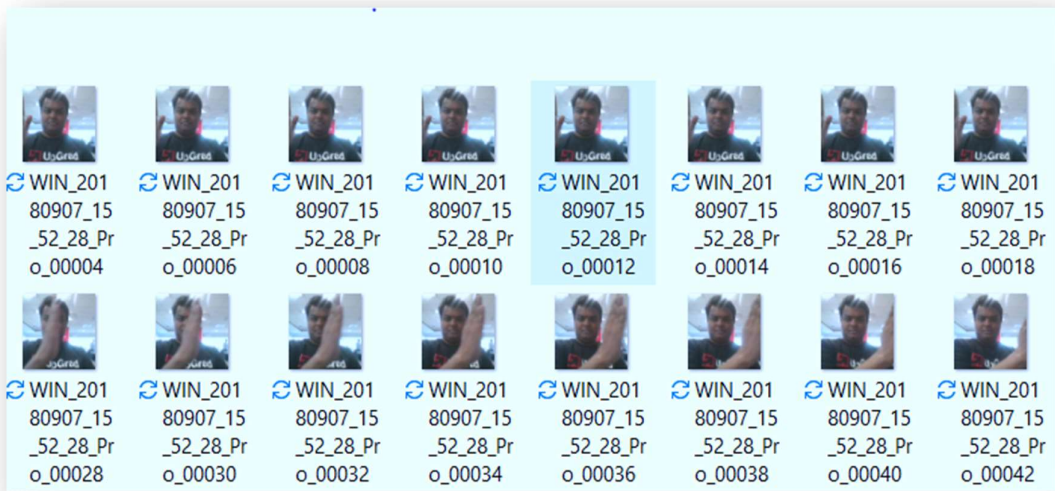
## Problem Statement

As a data scientist at a home electronics company which manufactures state of the art smart televisions. We want to develop a cool feature in the smart-TV that can recognise five different gestures performed by the user which will help users control the TV without using a remote.

- Thumbs up : Increase the volume.
- Thumbs down : Decrease the volume.
- Left swipe : 'Jump' backwards 10 seconds.
- Right swipe : 'Jump' forward 10 seconds.
- Stop : Pause the movie.

## Understanding the Dataset

The training data consists of a few hundred videos categorized into one of the five classes. Each video (typically 2-3 seconds long) is divided into a **sequence of 30 frames (images)**. These videos have been recorded by various people performing one of the five gestures in front of a webcam - similar to what the smart TV will use.



## Objective

Our task is to train different models on the 'train' folder to predict the action performed in each sequence or video and which performs well on the 'val' folder as well. The final test folder for evaluation is withheld - final model's performance will be tested on the 'test' set.

## Two types of architectures suggested for analysing videos using deep learning:

### 1. 3D Convolutional Neural Networks (Conv3D)

*3D convolutions* are a natural extension to the 2D convolutions you are already familiar with. Just like in 2D conv, you move the filter in two directions ( $x$  and  $y$ ), in 3D conv, you move the filter in three directions ( $x$ ,  $y$  and  $z$ ). In this case, the input to a 3D conv is a video (which is a sequence of 30 RGB images). If we assume that the shape of each image is  $100 \times 100 \times 3$ , for example, the video becomes a 4D tensor of shape  $100 \times 100 \times 3 \times 30$  which can be written as  $(100 \times 100 \times 30) \times 3$  where 3 is the number of channels. Hence, deriving the analogy from 2D convolutions where a 2D kernel/filter (a square filter) is represented as  $(f \times f) \times c$  where  $f$  is filter size and  $c$  is the number of channels, a 3D kernel/filter (a 'cubic' filter) is represented as  $(f \times f \times f) \times c$  (here  $c = 3$  since the input images have three channels). This cubic filter will now '3D-convolve' on each of the three channels of the  $(100 \times 100 \times 30)$  tensor

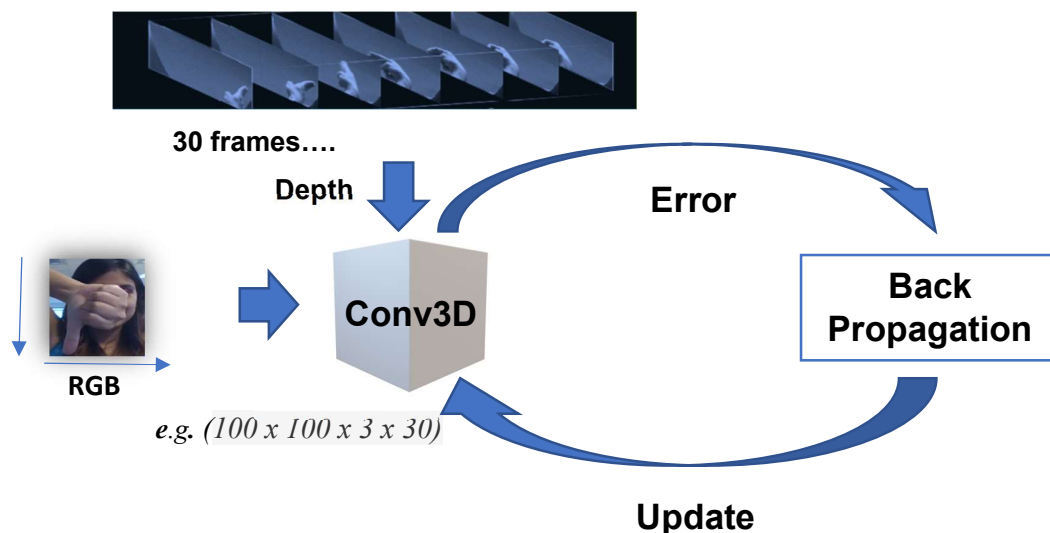


Figure 1: A simple representation of working of a 3D-CNN

### 2. CNN + RNN architecture

The *conv2D* network will extract a feature vector for each image, and a sequence of these feature vectors is then fed to an RNN-based network. The output of the RNN is a regular softmax (for a classification problem such as this one).





Figure 2: A simple representation of an ensembled CNN+LSTM Architecture

## Data Generator

This is one of the most important part of the code. In the generator, we are going to pre-process the images as we have images of 2 different dimensions ( $360 \times 360$  and  $120 \times 160$ ) as well as create a batch of video frames. The generator should be able to take a batch of videos as input without any error. Steps like cropping, resizing and normalization should be performed successfully.

## Data Pre-processing

- **Resizing and cropping of the images.** This was mainly done to ensure that the NN only recognizes the gestures effectively rather than focusing on the other background noise present in the image.
- **Normalization of the images.** Normalizing the RGB values of an image can at times be a simple and effective way to get rid of distortions caused by lights and shadows in an image.
- At the later stages for improving the model's accuracy, we have also made use of **data augmentation**, where we have slightly rotated the pre-processed images of the gestures in order to bring in more data for the model to train on and to make it more generalizable in nature as sometimes the positioning of the hand won't necessarily be within the camera frame always.



**CAUTION:** It was taken into consideration that we don't rotate images to a greater extent as this would change the meaning of the gestures completely 😊 !!

## NN Architecture development and training

- Experimented with different model configurations and hyper-parameters and various iterations and combinations of batch sizes, image dimensions, filter sizes, padding and stride length were experimented with. We also played around with different learning rates and *ReduceLROnPlateau* was used to decrease the learning rate if the monitored metrics (*val\_loss*) remains unchanged in between epochs.
- We experimented with *SGD()* and *Adam()* optimizers but went forward with *Adam()* as it lead to improvement in model's accuracy by rectifying high variance in the model's parameters. We were unsupportive of experimenting with *Adagrad()* and *Adadelata()* due to the limited computational capacity as these take a lot of time to converge because of their dynamic learning rate functionalities.
- We also made use of *Batch Normalization*, *pooling* and *dropout layers* when our model started to overfit, this could be easily witnessed when our model started giving poor validation accuracy inspite of having good training accuracy ☺.
- *Early stopping* was used to put a halt at the training process when the *val\_loss* would start to saturate / model's performance would stop improving.

## Observations

- It was observed that as the Number of trainable parameters increase, the model takes much more time for training.
- **Batch size  $\propto$  GPU memory / available compute.** A large batch size can throw *GPU Out of memory error*, and thus here we had to play around with the batch size till we were able to arrive at an optimal value of the batch size which our GPU could support
- Increasing the batch size greatly reduces the training time but this also has a negative impact on the model accuracy. This helped realise that there is always a trade-off here on basis of priority -> If we want our model to be ready in a shorter time span, choose larger batch size else you should choose lower batch size if you want your model to be more accurate.
- *Data Augmentation* and *Early stopping* greatly helped in overcoming the problem of overfitting which our initial version of model was facing.
- *CNN+LSTM* based model with *GRU* cells had better performance than *Conv3D*. As per our understanding, this is something which depends on the kind of data we used, the architecture we developed and the hyper-parameters we chose.
- *Transfer learning* **boosted** the overall accuracy of the model. We made use of the [\*MobileNet\*](#) Architecture due to it's light weight design and high speed performance coupled with low maintenance as compared to other well-known architectures like VGG16, AlexNet, GoogleNet etc.
- For detailed information on the Observations and Inference, please refer Table 1.

MODEL	EXPERIMENT	RESULT	DECISION + EXPLANATION	PARAMETERS
Conv3D	1	OOM Error	Reduce the batch size and Reduce the number of neurons in Dense layer	-
	2	Training Accuracy : 0.99 Validation Accuracy : 0.81	Overfitting Let's add some Dropout Layers	1,117,061
	3	Training Accuracy : 0.65 Validation Accuracy : 0.52 (Best weight Accuracy, Epoch:6/25)	Val_loss didn't improve from 1.24219 so early stopping stop the training process. Let's lower the learning rate to 0.0002.	3,638,981
	4	Training Accuracy : 0.76 Validation Accuracy : 0.72 (Best weight Accuracy, Epoch:12/25)	Overfitting has reduced but accuracy hasn't improved. <i>Let's trying adding more layers</i>	1,762,613
	5	Training Accuracy : 0.83 Validation Accuracy : 0.76	<i>Don't see much performance improvement. Let's try adding dropouts.</i>	2,556,533
	6	Training Accuracy : 0.84 Validation Accuracy : 0.69	Overfitting Increase, adding dropouts has further reduced validation accuracy. Let's try to reduce the parameters	2,556,533
	7	Training Accuracy : 0.84 Validation Accuracy : 0.74	Overfitting reduced, but validation accuracy low. Let's try to reduce the parameters. Val Accuracy: 0.49, Train Accuracy: 0.54	696,645
	8	Training Accuracy : 0.82 Validation Accuracy : 0.73	Accuracy remains below same. Let's switch to CNN+LSTM.	504,709
CNN+LSTM	9 (Model-8 on Notebook)	Training Accuracy : 0.93 Validation Accuracy : 0.85	CNN - LSTM model - we get a best validation accuracy of 85%.	1,657,445
Conv3D	Let's apply some <b>Data Augmentation</b> techniques & check the model performance			
	10	Training Accuracy : 0.78 Validation Accuracy : 0.82	(3, 3, 3) Filter & 160 x 160 image resolution	3,638,981
	11	Training Accuracy : 0.72 Validation Accuracy : 0.75	(2, 2, 2) Filter & 120 x 120 image resolution. Increase epoch count to 20. Network is generalizing well.	1,762,613
	12	Training Accuracy : 0.87 Validation Accuracy : 0.78	Adding more layers.	2,556,533
	13	Training Accuracy : 0.65 Validation Accuracy : 0.25	Very low performance. Let's reduce the network parameters.	2,556,533
	14	Training Accuracy : 0.89 Validation Accuracy : 0.78	After reducing network parameters, model's performance is quite good.	696,645
	15	Training Accuracy : 0.88 Validation Accuracy : 0.81	Reducing network parameters again.	504,709
	16	Training Accuracy : 0.98 Validation Accuracy : 0.77	Overfitting is considerably high, not much improvement.	2,573,541
CNN LSTM with GRU	17	Training Accuracy : 0.85 Validation Accuracy : 0.58	<i>We are not training the MobileNet weights that can see, validation accuracy is very poor.</i>	3,840,453
Transfer Learning(Optional)	18	Training Accuracy : 0.98 Validation Accuracy : 0.93	Awesome result !!	3,692,869
Transfer Learning with GRU &(Optional)				

Table 1: Observations and Results for numerous tested NN architectures

After doing all the experiments, the model finalized **Model 8–CNN+LSTM**, which performed well.

**Reason:**

- (Training Accuracy: 93%, Validation Accuracy: 85%)
- Number of Parameters (1,657,445) less according to other models' performance
- Learning rate gradually decreasing after some Epochs

### Further suggestions for improvement:

- **Using Transfer Learning:** Using a pre-trained *ResNet50/ResNet152/Inception V3* to identify the initial feature vectors and passing them further to a *RNN* for sequence information before finally passing it to a softmax layer for classification of gestures. (This was attempted but other pre-trained models couldn't be tested due to lack of time and disk space in the [nimblebox.ai](https://nimblebox.ai) platform.)
- **Using GRU:** A *GRU* model in place of *LSTM* appears to be a good choice. Trainable Parameters of a *GRU* are far less than that of a *LSTM*. Therefore would have resulted in faster computations. However, its effect on the validation accuracies could be checked to determine if it is actually a good alternative over *LSTM*.
- **Deeper Understanding of Data:** The video clips were recorded in different backgrounds, lightings, persons and different cameras were used. Further exploration on the available images could give some more information about them and bring more diversity in the dataset. This added information can be exploited in favour inside the generator function adding more stability and accuracy to model.
- **Tuning hyperparameters:** Experimenting with other combinations of hyperparameters like, activation functions (*ReLU*, *Leaky ReLU*, *mish*, *tanh*, *sigmoid*), other optimizers like *Adagrad()* and *Adadelta()* can further help develop better and more accurate models. Experimenting with other combinations of hyperparameters like the *filter size*, *paddings*, *stride\_length*, *batch\_normalization*, *dropouts* etc. can further help improve performance.